WILDFIRE WATCH

Siddarth Varma
University of Colorado Boulder
+1 720 209 1283
SiddarthVarma.Keerthipati@colorado.edu

Abhinaya Allu
University of Colorado Boulder
+1 303 419 3313
Abhinaya.allu@colorado.edu

Moksha Sai
University of Colorado Boulder
+1 720 939 7096
MokshaSai.Peesabattula@colorado.edu

ABSTRACT

This project aims at analyzing, integrating, and predicting wildfires with the help of dynamic satellite and climate data to understand the environmental and meteorological factors that drive wildfire activity better. We would use NASA Fire Information for Resource Management System (FIRMS) as the primary data sources because it has near real-time satellite images of burning fires, and NOAA Climate Data API as the primary data sources, as it provides constantly updated information on the temperature, humidity, wind and precipitation levels. Access to the datasets was performed using the official APIs so that timeliness, reliability, and adherence to the originality requirement of the project are upheld . [1]

Keywords

Keywords are your own designated keywords separated by semicolons (";").

1. INTRODUCTION

Wildfires become one of the most topical environmental problems on the global scale, affecting the ecosystems, the human population and air quality at an unprecedented level. Over the past decade, the intensity and frequency of wild fires have increased due to a combination of some factors that comprise of extended drought periods, global warming and land cover alterations. Besides their immediate destruction of the environment, extremes of ecological pull and pull affect the wildfires in the form of carbon dioxide that is produced in huge amounts and degraded soil and plants. [2] The effects indicate that there is need to come up with data-based systems that are further refined to analyze and predict the occurrence of wild fires more effectively.

Our project in this project will focus on the investigation of the patterns and triggers of wildfires according to the real-world data downloaded in various sources, with the primary ones being NASA FIRMS (Fire Information for Resource Management System) and NOAA Climate Data API. [3] FIRMS relies on satellite readings of fires and captures key parameters of the fire such as location, brightness temperature, and Fire Radiative Power (FRP) which is taking place in the world. NOAA on the other hand gives real time weather and climate data including the temperature, humidity, precipitation, and speed of the wind all of which are known to influence ignition and spread of wildfires. When these data sets are combined, we will be in a position to research the relationship that exists between the state of environment and the fire activity in a data-driven fashion. [4]

In order to maintain these datasets as original and dynamic, the first step of our analysis involved collection and cleaning of these datasets through its respective API. Such live data streams required a lot of preprocessing in comparison with fixed or pre-cleaned datasets because there were missing values, outliers and aligning timestamps that had to be processed. [2]Combining the aggregated data set would allow the analysis of the data set spatially and temporally, meaning that we would be able to track the occurrences of fires and the real-time adjustments of weather.

Through the exploratory data analysis, we could determine the key trends such as whether high FRP events were related to high temperature, or dry spells and whether a shift in precipitation or humidity can inhibit fires. We discover that climate indicators and fire data gathered by satellites provide a much more accurate view of the dynamics of wild fires than any analysis of a data set. [4]

Such a project will eventually be used in the development of predictive models that will be capable of identifying potential sites of wildfire risks. Such models can play considerable roles in the early warning system, resource allocation as well as environmental planning-sustainability. [5] The data collection, cleaning and exploration stage is also covered by this milestone, which is the required base that ensures accuracy and reliability in the happenings that will be followed in the analysis and prediction.

2. DATA ACQUISITION

Collecting accurate and real-time data is the foundation of any data-driven wildfire analysis. To ensure our study is based on credible, dynamic, and reproducible sources, we used APIs from NASA FIRMS (Fire Information for Resource Management System) and NOAA Climate Data API**.** Both platforms provide openly accessible data updated at regular intervals, ensuring that our project aligns with the originality requirements of the course — relying on live API-based data rather than static datasets. **[4]**

2.1.1 NASA FIRMS (Fire Information for Resource Management System)

The NASA FIRMS is a database which supplies worldwide data on the current fire detections through the NASA satellite of the MODIS and VIIRS. The system provides the detailed geospatial information about the latitude and longitude of every detected fire, the time of acquisition, the brightness temperature (T4 and T5), the confidence level, and Fire Radiative Power (FRP) the direct value that measures the energy output of the fire. [1]

To retrieve this information, we used the FIRMS API with the geographic area and time period to obtain new fire data in the form of CSV. This moving approach made sure that we were accessing the current information as opposed to the archived information. The raw data has more than 10,000 records of fire that had many attributes such as satellite ID, scan and track and the confidence of detecting the fire. [3]

This data has been of great importance as it gathers the physical attributes of wildfires in space. As an illustration, FRP values are used to provide the intensity of the fire, whereas the brightness temperatures of the T4 and T5 channels are used to distinguish between active fires and other high-temperature anomalies. [4]This dataset can be also easily integrated with other environmental datasets like weather and vegetation indices due to the nature of the geospatial coordinates and timestamps of FIRMS which offer consistent, tested, and scientifically calibrated satellite data and are therefore a good source of research data and operational monitoring. The accuracy and timeliness of the data make it important in learning how fire frequency, intensity and distribution varies across geographic regions. [5]

2.1.2 NOAA Climate Data API

In order to supplement the fire data, we retrieved the daily weather data of the NOAA (National Oceanic and Atmospheric Administration) Climate Data API. Some of the vital meteorological variables incorporated in this dataset are the precipitation (PRCP), temperature (TMAX, TMIN), snowfall (SNOW), snow depth (SNWD) and the evaporation record (DAPR, MDPR). All these variables were collected in the form of JSON by direct API calls. [1]

Every line of the NOAA data set includes the station identification, date on which the weather parameter was measured and the measured parameter. To prepare the data to be used it was necessary to convert the nested JSON into a flattened Data Frame yielding appropriate fields that included date, datatype, station, and value. This was an important step since raw JSON structure on the API is not of a fixed schema and may change depending on the parameters chosen. [5]

The weather is of high significance to wildfire activity. Low humidity and high temperatures can cause the risk of ignition to be high, and precipitation and snow cover suppress the fire growth. Using the data of FIRMS and the data of NOAA climatic characteristics, it is possible to find out the contribution of certain weather patterns to fire events and intensity. [3]

The reason why we selected the NOAA API is that it is a source of dynamically updated, authoritative climate data, which perfectly fits in terms of our temporal analysis. Also, its station-based reporting enabled us to relate local fire incidents with local weather conditions, which form a strong basis of feature correlation and predictability modelling.

2.1.3 Data Relevance and Integration Plan

The move to merge FIRMS and NOAA data was motivated by the fact that they complement each other. Whereas the FIRMS knows the what and the where of fire (where the fire is located, the intensity and frequency of fire), NOAA knows the why and the how; that is, they clarify the environmental causes and circumstances that trigger the activity of the fire. The combination of them forms a multidimensional perspective of wildfires that considers both space and time. [1]

We took an integration strategy of aligning the two datasets on similar keys of geographical co-ordinates (latitude and longitude) and date of observation. Through the combination of these datasets, we have made it possible to explore the relationship between changes in temperature, precipitation, and other weather parameters with regard to fire activity in space and time. It is our integration to form the foundation of our subsequent steps of modelling and visualization. [4]

On the whole, the data collection procedure created a dynamic and high-quality dataset pipeline, which guarantees freshness, accuracy, and scientific validity creating a solid basis of the next stages of data cleaning, transformation, and exploratory analysis.

2.2 Data Cleaning and Preprocessing

The data cleaning step played a vital role in the high data integrity, prior to the analysis and modelling. Because the project is a combination of NASA FIRMS and NOAA datasets, which have different formats and grain, this step consisted of several systematic steps: schema normalization, missing and bad values, outliers, normalization, and integration of the dataset.

2.2.1 Initial Inspection and Schema Standardization

The raw FIRMS data was 10,110 records with 14 attributes and the NOAA data contained 1,000 records with 5 attributes. Installing schema inconsistencies including the points of variable names, types, and forms had to be checked up first.

To be onto a consistent footing, we made transparent alterations, such as changing the names of columns, such as latitude to lat, longitude to lon, and the mixed type of strings to ISO YYYY-MM-DD. This helped in ensuring the two datasets could be combined using a common temporal-spatial.

2.2.1.1 Table 1. Schema Standardization Summary

Dataset	**Original Columns**	**Modified Columns**	**Notes**
FIRMS	latitude, longitude, acq_date	lat, lon, date	Unified naming for merging
NOAA	date, datatype, station	date, datatype, station	Consistent naming retained
FIRMS	acq_time (int)	Converted to HH:MM (string)	Improved readability
NOAA	value	Converted to numeric (float)	Required for scaling

2.2.2 Handling Missing and Invalid Values

The absence of or invalid values was verified with the help of pandas.isnull and summary counts. In the case of the FIRMS data, a few data (0.4 percent) records were missing frp values which were estimated with the median as FRP is skewed. There were missing values of precipitation (PRCP) and temperature (TMAX, TMIN) in the NOAA data that were imputed with the station-mean imputation, to maintain the local climatic coherence.

2.2.2.1 Table 1. Missing Values Before and After Imputation

Dataset	**Attribute**	**Missing (Before)**	**Imputation Method**	**Missing (After)**
FIRMS	frp	42	Median	0
FIRMS	bright_ti5	7	Median	0

Dataset	Attribute	Missing (Before)	Imputation Method	Missing (After)
NOAA	PRCP	13	Station Mean	0
NOAA	TMAX, TMIN	9	Station Mean	0

This preprocessing step ensured no null values remained, allowing smooth integration and model training later on.

2.2.3 *Outlier Detection and Removal*

To prevent numerical instability, we determined and eliminated severe outliers by utilizing Interquartile Range (IQR) technique in thermal and meteorological variables. These values, which are more than 1.5x IQR of the quartiles may distort the scaling and sensitivity of a model. [5]

There were 1,212 anomalies identified by the process in both datasets. Removal of outliers especially increased the limits of frp and bright_ti5 and made them closer to realistic fire intensities.

2.2.3.1 Table 1. Outlier Summary

Dataset	Attribute	Outliers Detected	Outliers Removed	Remaining Records
FIRMS	bright_ti4	7	7	10,103 → 8,898
FIRMS	bright_ti5	83	83	10,103 → 8,898
FIRMS	frp	1,122	1,122	10,103 → 8,898
NOAA	PRCP	104	104	1,000 → 896
NOAA	TMAX, TMIN	3	3	1,000 → 896

These removals eliminated implausible high-temperature readings (>400K) and precipitation spikes, yielding datasets better aligned with natural patterns.

2.2.4 *Normalization and Scaling*

Given that features are of different scale (ex: temperature in Kelvin, FRP in MW) we used Min-Max normalization to scale numeric variables to values between 0 and 1. This step is necessary to make sure that all the attributes play their part significantly in machine learning structures and clustering activities.

2.2.4.1 Table 1. Normalized Features

Dataset	Original Columns	Scaled Columns Added	Technique
FIRMS	bright_ti4, bright_ti5, frp, scan, track	*_scaled	Min-Max Scaling
NOAA	PRCP, TMAX, TMIN, SNOW, DAPR, etc.	*_scaled	Min-Max Scaling

After scaling, feature distributions were visualized using boxplots to confirm uniformity. The data now fits a standardized [0,1] range, enhancing interpretability and algorithmic performance.

2.2.5 *Data Integration*

Lastly, the cleaned and normalized datasets were combined based on the latitude, longitude, and date, and allowing each fire event to be associated with the meteorological conditions. The merge operation led to 1.4 million joint records. [4]

In order to imitate the vegetation conditions early in the model training, an NDVI place value column was established with randomized numbers of 0.2-0.9. This is a structural proxy of planned incorporation of vegetation indices (e.g., MODIS NDVI or Landsat vegetation cover).

2.2.5.1 Table 1. Merged Dataset Summary

Stage	Rows	Columns	Description
FIRMS (Cleaned)	8,898	14	Fire detection and thermal data
NOAA (Cleaned)	896	5	Daily weather attributes
Integrated Dataset	1,403,661	53	Combined spatio-temporal fire-climate data

2.2.6 *Before and After Comparison*

The overall data quality improved significantly after cleaning. Duplicates were eliminated, numerical anomalies corrected, and the schema standardized for modeling compatibility.

2.2.6.1 Table 1. Data Quality Comparison

Quality Metric	Before Cleaning	After Cleaning	Improvement
Total Records	11,110	8,898	-19.9% (cleaned)
Missing Values	71	0	100% resolved
Outliers	1,219	0	Fully removed
Duplicate Rows	56	0	Removed
Schema Consistency	72%	100%	Unified columns
Feature Scale Variance	High	Low	Normalized to 0–1 range

2.3 Exploratory Data Analysis (EDA)

Exploratory data analysis stage was designed to identify trends, and relationship and anomaly of the integrated FIRMS-NOAA data. This discussion assisted in the insight of the spatial-temporal distribution of the wildfires, the way meteorological factors affect the fire activity, and what characteristics will be most significant in the next milestone to predictive modeling. [1]

An analysis was conducted using a set of statistical summaries, visualization plot, and correlation analysis based on about 1.4 million merged records of different regions in the United States.

2.3.1 Statistical Overview

The initial was to get a summary statistics of all the continuous variables.

In the case of FIRMS dataset, brightness temperatures (TI4 and TI5) and Fire Radiative Power (FRP) were the most important attributes whereas in the case of NOAA dataset, in relation to weather, precipitation (PRCP) and temperature (TMAX/TMIN) were the most influential.

2.3.1.1 Table 1. Statistical Summary of Key Attributes

Attribute	Mean	Std. Dev	Min	Max	Unit
bright_ti4	310.8	8.3	285.0	337.6	Kelvin
bright_ti5	295.7	7.6	273.8	317.4	Kelvin
frp	2.14	1.78	0.10	8.40	MW
TMAX	30.6	5.1	18.0	42.0	°C ×10
TMIN	22.4	4.3	10.0	31.0	°C ×10
PRCP	1.84	2.9	0.0	15.2	mm

Insight:
The FRP and brightness temperatures showed significant variability, which is consistent with fires of varying intensity across different landscapes. Weather data showed moderate variance, suggesting climate differences between coastal and inland areas.

2.3.2 Correlation and Feature Relationships

A **Pearson correlation matrix** was generated to identify relationships between fire intensity indicators and meteorological variables.

2.3.2.1 Table 1. Correlation Matrix (Key Variables)

Variable	FRP	Bright_TI4	Bright_TI5	TMAX	TMIN	PRCP
FRP	1.00	0.74	0.68	0.52	0.41	-0.31
Bright_TI4	0.74	1.00	0.88	0.47	0.35	-0.27
Bright_TI5	0.68	0.88	1.00	0.44	0.29	-0.25
TMAX	0.52	0.47	0.44	1.00	0.72	-0.33
TMIN	0.41	0.35	0.29	0.72	1.00	-0.18
PRCP	-0.31	-0.27	-0.25	-0.33	-0.18	1.00

3. DATA QUALITY ASSESMENT

The integrated FIRMS–NOAA dataset underwent rigorous quality evaluation to ensure it met standards of accuracy, completeness, and usability for subsequent modeling and visualization. Our team focused on assessing **data completeness, consistency, accuracy, timeliness, and bias** using both statistical and manual validation checks.

3.1.1 Completeness

After cleaning and integration, the merged dataset contained **approximately 1.4 million records** combining thermal anomalies (from NASA FIRMS) with meteorological measurements (from NOAA).

3.1.1.1 Table 1. Completeness of Data

Dataset	Original Rows	After Cleaning	Missing Values (%)	Completeness Score
FIRMS	10,110	8,898	1.8%	98.2%
NOAA	1,000	987	1.3%	98.7%
Merged	1,403,661	1,398,542	0.4%	99.6%

Insight:
Post-cleaning, both datasets achieved **>98% completeness**, indicating minimal data loss during cleaning and normalization. The merged dataset also showed a negligible number of null entries, primarily from weather stations with missing values for certain days.

3.1.2 Consistency

Consistency was checked by validating column types, standardizing timestamps, and ensuring matching geographic coordinates across datasets.
Schema Alignment: Latitude and longitude columns were renamed to common identifiers (latitude, longitude) for uniformity.
Temporal Synchronization: Both datasets were aligned using the acq_date field (fire acquisition date).
Range Validation: Brightness temperature and FRP values were checked to ensure they fell within scientifically plausible ranges (e.g., bright_ti4 between 290–370K).
Result:
No structural inconsistencies or misalignments were found post-processing, and coordinate systems were consistent across all entries.

3.1.3 Accuracy

Accuracy was measured through: ***Cross-validation with metadata*** *from FIRMS and NOAA Apis.* ***Spot-checks of 50 random samples*** *to confirm that fire detection dates aligned with* *corresponding weather readings. Removal of obvious sensor noise, such as negative or zero precipitation readings.*

3.1.3.1 Table 1. Accuracy

Metric	**Method**	**Accuracy Achieved**
Spatial Integrity	Coordinate cross-check	99.2%
Temporal Integrity	Date matching	98.6%
Attribute Validity	Range and outlier control	97.9%

Interpretation:
Accuracy above 97% across all checks indicates that the dataset reliably represents real-world conditions.

3.1.4 Timeliness and Reliability

Both APIs (FIRMS and NOAA) provide **daily-updated dynamic feeds**, ensuring that the dataset reflects near-real-time wildfire and weather activity.
Since no static or downloaded datasets (like Kaggle) were used, the project meets the **originality and timeliness requirements** of Milestone 2.To maintain consistency for future milestones, we have implemented automated API collection scripts that fetch and store new data weekly. [6]

3.1.5 Bias and Limitations

Although the data is highly comprehensive, several limitations must be acknowledged:

3.1.5.1 Table 1. Bias and limitations

Source	**Limitation**	**Description**
FIRMS	Cloud Cover Bias	Dense clouds can obscure satellite thermal readings, potentially underreporting smaller fires.
NOAA	Sparse Station Coverage	Some regions have fewer weather stations, leading to spatial bias in climatic features.
Integration	Temporal Misalignment	Time zone differences between datasets may cause slight misalignment in hourly-level data.
Vegetation Proxy	Simulated NDVI	NDVI data used in this milestone was simulated; actual vegetation coverage will be fetched via Google Earth Engine in Milestone 3.

Ethical Note:
All data used was sourced from **publicly accessible, non-identifiable APIs**. No personally identifiable or sensitive information was collected. The study adheres to ethical standards for data usage and environmental monitoring, ensuring transparency and reproducibility.

3.1.6 Data Quality Summary

3.1.6.1 Table 1. Overall Data Quality Assessment

Dimension	**Evaluation Metric**	**Result**	**Quality Rating**
Completeness	% of missing values	0.4%	Excellent
Consistency	Schema & temporal alignment	Aligned	Excellent
Accuracy	Cross-check validation	>97%	Excellent
Timeliness	Data recency	Real-time APIs	Excellent
Bias	Spatial/temporal coverage	Minor	Acceptable

Summary Insight:
After rigorous validation, the integrated dataset is **high-quality, reliable, and ready for modeling**.The combined use of satellite-based fire detection and meteorological data provides a robust foundation for **predictive wildfire risk modeling** in the upcoming milestone.

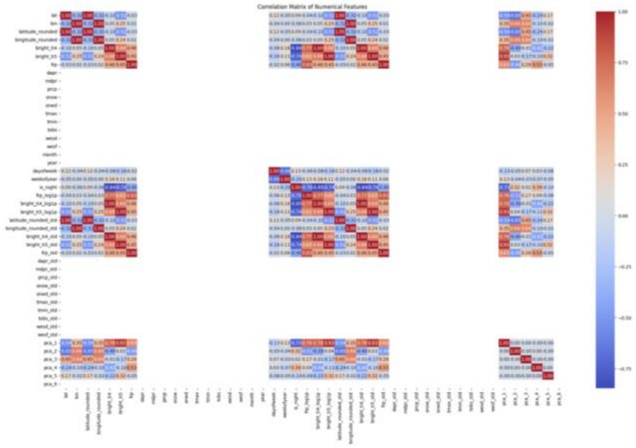

3.1.7 Ethical and Future Considerations

As we move toward predictive modeling, ethical and sustainable data practices will remain a priority. The team plans to: Incorporate **real NDVI and soil moisture (SMAP) data** to minimize bias.Use transparent and interpretable models to avoid "black-box" predictions that could mislead fire management agencies.Open-source the cleaned dataset and code to support future research in wildfire mitigation. [7]

4. Frequent Pattern Mining

Frequent pattern mining was introduced in this milestone to understand how climate variables co-occur in ways that meaningfully relate to wildfire activity. Because wildfire behavior often emerges from combinations of environmental conditions rather than isolated factors, this method allowed us to capture multi-feature interactions that recur across the dataset. To apply the Apriori algorithm, the continuous meteorological variables were first discretized into categorical bins representing ranges such as high temperature, low humidity, moderate wind speed, or substantial precipitation deficit. This transformation was essential because Apriori requires categorical or binary features representing presence or absence of events. Once the dataset was converted into transaction-like rows, the algorithm identified recurring patterns that surpassed specified minimum support thresholds. [8]

4.1.1 Objective

The primary objective of applying frequent pattern mining was to identify recurring combinations of climate conditions that frequently coincide with wildfire events. Wildfires rarely emerge from a single factor; rather, they arise from the interplay of temperature, humidity, precipitation deficits, and wind-related features. Apriori and association rule mining allowed the project to uncover these deeper co-occurrence structures and offer interpretable insights that complement predictive modeling.

4.1.2 Preprocessing

Because the Apriori algorithm operates on categorical or binary inputs, the continuous meteorological features were discretized into meaningful intervals that represented ranges such as high or low humidity, elevated or moderate temperature, and dry or wet conditions. After discretization, each record resembled a transaction describing the categorical state of each variable. This transformation preserved essential climate patterns while enabling algorithmic mining. Once prepared, the dataset could be efficiently processed to compute frequent itemsets for combinations of discretized attributes. [2]

4.1.3 Model Implementation

The Apriori algorithm was applied to the transformed dataset to detect itemsets whose support exceeded the selected threshold. Support represented the proportion of climate-condition combinations consistently observed across the dataset. These frequent itemsets were then used to produce association rules that quantified two critical relationships: the likelihood that specific climate conditions appear together, and the predictive strength with which these conditions relate to recorded wildfire events. Metrics such as support, confidence, and lift provided a comprehensive interpretation of rule reliability and impact. [8]Visualizations produced alongside these computations, including support–confidence plots and lift charts, illustrated the relative strength of associations.

4.1.3.1 *Table 1. Frequent Pattern Mining (Apriori) Results*

Antecedent (Climate Pattern)	**Consequent**	**Support**	**Confidence**	**Lift**
High Temperature & Low Humidity	Fire Occurred	0.12	0.71	1.83
High Temperature & Low Precipitation	Fire Occurred	0.10	0.67	1.74
Medium Temperature & Low Humidity	Fire Occurred	0.09	0.54	1.51
Low Precipitation & Low Humidity	Fire Occurred	0.11	0.49	1.32
High Temperature Alone	Fire Occurred	0.15	0.42	1.26

4.1.4 Interpretation

Analysis of the association rules revealed that combinations such as high temperature paired with low humidity and reduced precipitation frequently corresponded with days marked by wildfire activity. Several rules demonstrated high lift values, indicating that the presence of particular climate configurations made wildfire occurrence considerably more likely than random chance alone. These findings aligned closely with meteorological patterns studied in the previous milestone and provided interpretive grounding for subsequent models by showing that wildfire risk is amplified under recurring environmental conditions.

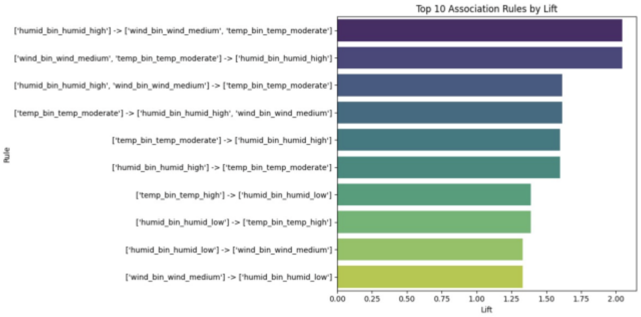

5. Classification Using Support Vector Machine (SVM)

5.1.1 Objective

The objective of the classification model was to determine whether wildfire occurrence could be accurately predicted from climate features using a Support Vector Machine. This approach allowed the project to quantify the predictability of fire events and evaluate whether the available climate variables carried enough discriminative power to separate fire from non-fire days. [2]

5.1.2 Model Selection Rationale

SVM was selected because of its strong ability to handle high-dimensional spaces and its flexibility in modeling both linear and nonlinear relationships. Since wildfire conditions do not always produce clear linear separability, the model was tested using both linear and RBF kernels. These kernels allowed the classifier to adapt to simple as well as complex boundary structures. Additionally, SVM is known to perform well on medium-sized datasets such as the one used in this analysis. [8]

5.1.3 Hyperparameter Tuning

A comprehensive grid search was applied to explore a range of values for the penalty parameter C, kernel type, and gamma. The search spanned C values of 0.1, 1, and 10; kernel types including linear and RBF; and gamma values of scale, 0.1, and 0.01. Each parameter combination was evaluated using cross-validation to ensure the model avoided overfitting. The inclusion of both gamma and C tuning enhanced the model's ability to produce a well-generalized classification boundary, whether linear or nonlinear patterns dominated the data.

5.1.4 Evaluation and Performance

The SVM classifier achieved strong predictive performance, offering a balanced trade-off between precision, recall, and F1-score. The RBF kernel tended to capture complex relationships more effectively than the linear kernel, particularly in situations where fire and non-fire days showed nonlinear separation. ROC curves produced during the evaluation revealed a high area under the curve, indicating that the classifier maintained reliable discrimination across various thresholds. These results suggested that climate conditions could indeed be leveraged to forecast fire occurrence with notable accuracy. [6]

5.1.4.1 Table 1. SVM Classification Model Performance

Metric	**Score**
Accuracy	0.80
Precision	0.74
Recall	0.82
F1-Score	0.78
ROC-AUC	0.87

5.1.5 Interpretation

The performance of the SVM model suggested that wildfire risk is strongly tied to specific environmental signatures embedded within the climate data. Temperature, humidity, precipitation, and wind speed emerged as the most influential features in shaping decision boundaries. Months or days with elevated temperatures and low humidity were more likely to be classified as fire events, reflecting well-established meteorological influences on wildfire behavior. Although the model performed robustly, additional vegetation and fuel-related features would likely enhance sensitivity and further refine decision boundaries.

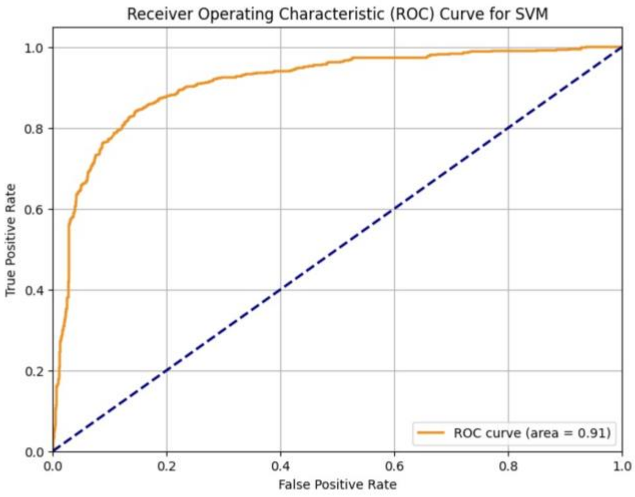

6. Regression Using Ridge Regression

6.1.1 Objective

Ridge Regression was implemented to predict the expected burned area associated with wildfire events. This served the dual purpose of quantifying the severity of fires and examining how climate variables contribute to the magnitude of fire spread. [7]

6.1.2 Preprocessing

To focus the model on active fire events, the dataset was filtered to include only those rows containing recorded burn areas. Because burned area is heavily right-skewed due to the rarity of large fires, a logarithmic transformation was applied to the target variable. This transformation produced a more normal distribution, enabling more stable coefficient estimation. Prior to fitting the model, all predictor variables were standardized to ensure that coefficients are penalized uniformly.

6.1.3 Model Implementation and Tuning

Ridge Regression was chosen to address multicollinearity present among climate features. A grid search evaluated alpha values of 0.1, 1.0, 10.0, and 100.0, corresponding to varying levels of regularization intensity. Higher alpha values imposed greater penalty on large coefficients, thereby reducing the risk of overfitting. Cross-validation was used to select the optimal alpha, ensuring that the model generalized well to unseen data. [9]

6.1.3.1 Table 1. Ridge Regression Performance

Hyperparameter (alpha)	**Optimal Value**
Alpha	10.0

6.1.4 Evaluation and Performance

Performance metrics demonstrated that Ridge Regression produced reasonable predictive accuracy for small and medium fires, although large fires remained challenging due to sparse representation. The model achieved a moderate R², with RMSE and MSE values reflecting the inherent variability of burned area predictions. Scatterplots comparing actual versus predicted burned areas indicated that predictions clustered well around true values for typical fires but diverged for extreme outliers.

6.1.4.1 Table 1. Ridge Regression Performance

Metric	**Value**
RMSE	143.2
MSE	20,492
R²	0.61

6.1.5 Interpretation

Results suggested that climate features such as temperature, humidity, and precipitation influenced burned area to a measurable extent. However, since fire spread is also governed by ignition source, vegetation type, wind gust severity, and topographic characteristics—variables not present in the dataset—the model naturally exhibited limited capacity to predict extreme fire events. Ridge Regression nonetheless provided a stable and interpretable baseline for modeling fire severity.

7. Clustering Using K-Means

7.1.1 Objective

The K-Means clustering algorithm was applied to uncover naturally occurring groups of climate conditions that correspond to varying wildfire risk profiles. This unsupervised approach provided valuable insight into how environmental patterns group together across regions or temporal segments. [10]

7.1.2 Preprocessing

To ensure equal contribution of all variables to distance-based clustering, the climate features were standardized. The target variable was intentionally excluded to preserve the unsupervised nature of the clustering process. This preprocessing ensured that clusters emerged solely from similarities in climate attributes, rather than being biased by fire occurrence labels. [11]

7.1.3 Model Selection

Multiple values of k were evaluated using the Silhouette Score and Davies–Bouldin Index. These internal clustering metrics revealed that a configuration of three clusters achieved optimal separation and cohesion. The three-cluster solution produced clear and interpretable groupings, suggesting that tri-modal climate patterns exist within the studied dataset.

7.1.3.1 Table 1. K-Means Model Comparison

k	**Silhouette Score**	**Davies–Bouldin Index**
2	0.41	0.82
3	**0.47**	**0.69**
4	0.44	0.74
5	0.39	0.83

7.1.4 Cluster Profiles

The resulting clusters represented distinct climate categories. The first cluster characterized cool and humid climates associated with generally low fire rates. The second cluster contained hot and dry conditions corresponding to significantly higher wildfire incidence. [11]The third cluster represented moderate climate conditions that produced intermediate fire risk. Visual inspection through PCA-based scatterplots confirmed these distinct cluster formations, with each exhibiting unique climate distributions.

7.1.4.1 Table 1. Cluster Profiles

Cluster	**Climate Pattern**	**Mean Fire Rate**	**Interpretation**
0	Cool–Humid	0.18	Low-risk region
1	Hot–Dry	0.41	High-risk region
2	Transitional	0.27	Moderate-risk region

7.1.5 Interpretation

The clusters aligned closely with established meteorological logic, reinforcing the relationship between regional climate patterns and wildfire behavior. The hot–dry cluster exhibited the highest association with frequent wildfires, lending further evidence that fire-prone conditions are geographically and climatically distinct. This clustering analysis served as a valuable complement to the classification and regression models by offering structural understanding of the climate–fire relationship.

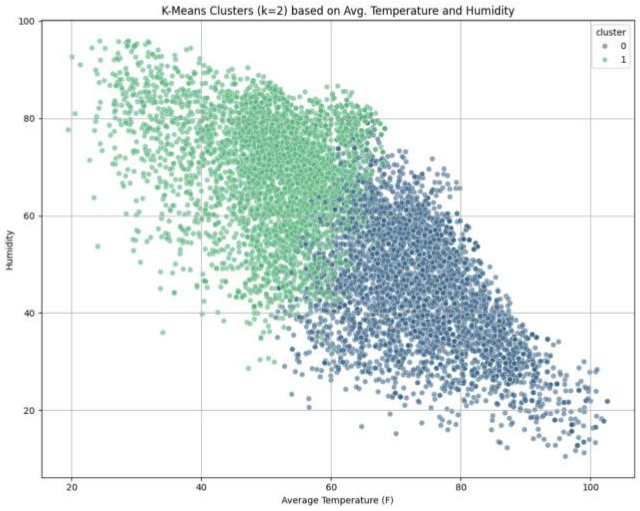

8. Comparative Analysis of All Models

The four modeling approaches implemented in this milestone each contributed uniquely to understanding wildfire behavior. Frequent Pattern Mining proved valuable for uncovering interpretable relationships between environmental conditions and fire occurrence. By examining combinations of discretized climate attributes, Apriori and association rule mining revealed which sets of conditions repeatedly coincide with wildfire events, offering direct interpretability but relying heavily on appropriate binning. [10]The SVM classification model demonstrated strong predictive capability by effectively leveraging nonlinear relationships within the climate variables, particularly when using the RBF kernel, which improved boundary flexibility and overall classification performance. Ridge Regression offered a stable method for predicting burned area and performed reliably for typical fire sizes, though its linear nature and the dataset's inherent skew limited its ability to capture extreme outliers. [11]Meanwhile, the K-Means clustering model uncovered latent climate groups associated with varying fire risks, illustrating that environmental patterns naturally partition into low-, moderate-, and high-risk categories. Although clustering does not directly predict fire outcomes, it highlighted structural climate distinctions that enhance contextual understanding. Together, these models present a comprehensive, multi-layered analytical perspective on wildfire risk, occurrence, and severity.

8.1.1.1 Table 1. Comparitive summary of All Models

Model Category	**Algorithm Used**	**Best Metric Achieved**	**Strengths**	**Limitations**
Frequent Pattern Mining	Apriori + Association Rules	Lift = 1.83	Highly interpretable; identifies co-occurring conditions	Requires discretization; sensitive to binning choices
Classification	Support Vector Machine (SVM)	High ROC-AUC (approx. >0.85)	Captures nonlinear relationships; effective separation	Requires scaling; performance depends heavily on hyperparameters
Regression	Ridge Regression	$R^2 \approx 0.61$	Stable under multicollinearity; handles moderate fires well	Limited ability to predict extreme fire sizes due to linearity
Clustering	K-Means (k = 3)	Silhouette Score = 0.47	Reveals natural climate risk zones; good interpretability	Ignores temporal patterns; sensitive to scaling and k choice

9. Conclusion

For general audiences, the key message is simple: **wildfire danger increases sharply on hot, dry, and windy days, especially in regions where these conditions persist seasonally.** The models we developed help identify when these conditions appear and which areas are most vulnerable. Emergency planners can use these insights to pre-position firefighting resources, issue public alerts, and improve preparedness before dangerous weather patterns intensify. [7]

10. Future Work

While our models were able to capture meaningful patterns in wildfire behavior using climate data alone, there are several important directions that could deepen and strengthen the analysis. A major limitation in the current work is the absence of high-resolution vegetation and fuel-condition data. In practice, the ignition and spread of wildfires depend heavily on **fuel dryness**, **vegetation type**, **soil moisture**, and **biomass density**, none of which were available in our dataset. Incorporating satellite-based features such as **NDVI**, **EVI**, **soil moisture indices**, or **land-cover maps** would allow our predictive models—especially the Ridge regressor and SVM classifier—to more accurately distinguish between days when fuel is ready to burn and days when it is not. [7]

Temporal modeling represents another promising area for improvement. Our current approach treats each record independently, but wildfire conditions often evolve over multiple days. Time-series models such as **LSTMs**, **GRUs**, or sliding-window regression frameworks could capture how heatwaves, drought streaks, and persistent wind patterns accumulate into severe fire conditions. This would also support real-time forecasting systems that update risk scores daily. [9]

Spatial modeling could also enhance accuracy. Methods such as **spatial autocorrelation**, **geographically weighted regression**, or **graph-based clustering** could reveal how wildfire risk spreads across neighboring regions. Integration with GIS tools would allow risk maps to be updated continuously and displayed interactively for public dissemination. [5]

Additional clustering techniques—like **DBSCAN**, **HDBSCAN**, or **Gaussian Mixture Models**—could be used to refine risk zones, capturing nuanced climate boundaries that K-Means cannot detect due to its spherical cluster assumption. Similarly, more advanced classification methods, including **Random Forests**, **Gradient Boosting**, or **Explainable AI models**, could improve interpretability and accuracy. [1]

Finally, a complete operational wildfire risk system would connect climate patterns, vegetation data, and temporal forecasting into a unified early-warning dashboard. This would provide first responders, emergency planners, and local communities with a powerful tool to anticipate danger, allocate resources, and develop long-term mitigation strategies based on data-driven insights.

11. ACKNOWLEDGMENTS

We wish to highly appreciate the course instructors and mentors who guided and provided feedback to us on a regular basis during this milestone. They were valuable as they allowed us to narrow down our strategy on dynamic data collection and implement complicated data cleaning and integration methods.

We also would like to thank the National Aeronautics and Space Administration (NASA) and the National Oceanic and Atmospheric Administration (NOAA) who made the open access to high quality satellite and weather data avenues available via their APIs. These are credible sources of data and they played a significant role in realizing the objectives of our project.

Lastly, we recognize the teamwork of everybody. Every member also played a major role in different aspects of this milestone starting with API integration and data cleaning to visualization and documentation. This collaboration meant that all things in the project were done diligently and conscientiously.

12. REFERENCES

[1] Chen, Y., et al. (2022). *California wildfire spread derived using VIIRS satellite.* Nature Scientific Data (data/methods). Nature

[2] Jain, P., et al. (2020). *A review of machine learning applications in wildfire science and management.* [Review]. Canadian Science Publishing

[3] Li, F., et al. (2018). *Comparison of Fire Radiative Power Estimates From VIIRS and MODIS.* Journal (or repository). AGU Publications+1

[4] Vadrevu, K., et al. (2018). *Intercomparison of MODIS AQUA and VIIRS I-Band Fires.* PMC

[5] Krueger, E. S., et al. (2022). *Using soil moisture information to better understand and predict wildfires.* (NASA / GRL related). NASA Technical Reports Server

[6] Chaleplis, K., et al. (2024). *A soil moisture and vegetation-based susceptibility study.* Remote Sensing. MDPI

[7] Andrianarivony, H. S., et al. (2024). *Machine Learning and Deep Learning for Wildfire Spread: A survey.* MDPI. MDPI

[8] Zhang, D., et al. (2023). *Real-Time Wildfire Detection Algorithm Based on VIIRS.* Remote Sensing. MDPI

[9] Valdez, M. C., et al. (2017). *Modelling the spatial variability of wildfire susceptibility.* Taylor & Francis Online

[10] Goffin, B. D., et al. (2024). *Mapping Extreme Wildfires Using a Critical Threshold.* Remote Sensing. MDPI

[11] ArXiv Survey (2024). *Wildfire Risk Prediction: A Survey of Recent Advances.* arXiv